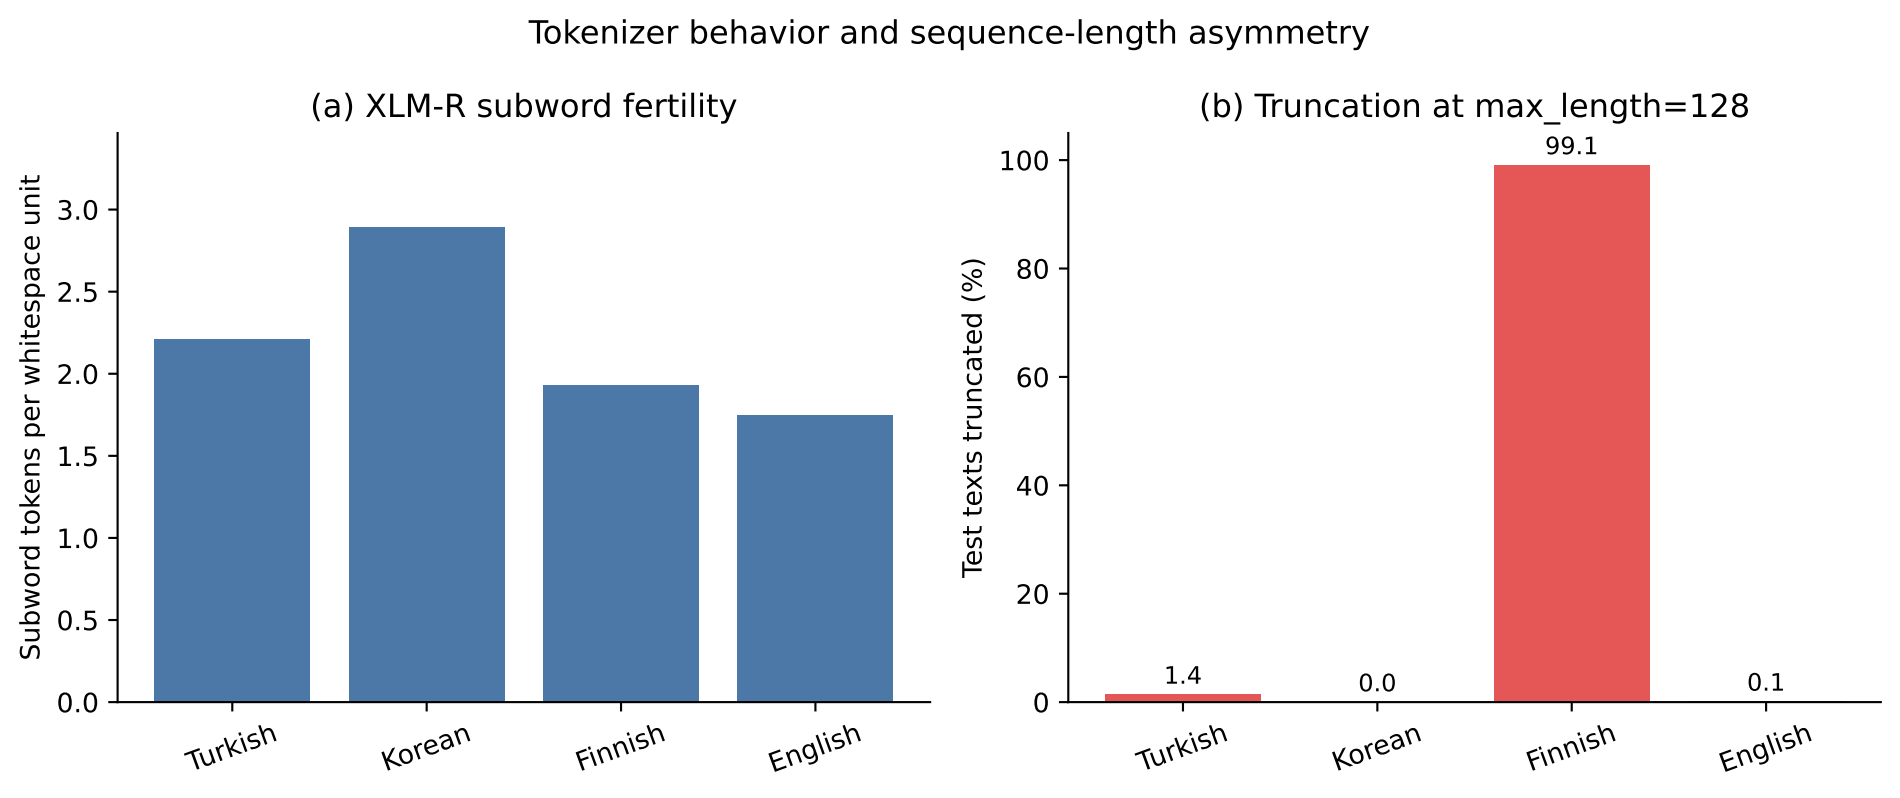
Tokenizer behavior and sequence-length asymmetry
(a) XLM-R subword fertility
Subword tokens per whitespace unit
0.0
0.5
1.0
1.5
2.0
2.5
3.0
Turkish
Korean
Finnish
English
(b) Truncation at max_length=128
Test texts truncated (%)
0
20
40
60
80
100
1.4
0.0
99.1
0.1
Turkish
Korean
Finnish
English